

HIVE Case Study – BI Track

Submitted by:
Prashant Kapoor
prashantkap1997@gmail.com

Problem Statement

We were asked to work with a public clickstream dataset of a cosmetics store and were given the task of extracting valuable insights from it using HIVE.

Here are the links for our datasets:

<https://e-commerce-events-ml.s3.amazonaws.com/2019-Oct.csv>

<https://e-commerce-events-ml.s3.amazonaws.com/2019-Nov.csv>

Note: In this report, all the steps taken have been outlined along with the screenshots. All the commands in text format are present at the end in the appendix section.

Solution/Implementation

Launching EMR and loading Data on Hive

emr 5.29.0 cluster was launched with 1 master and 1 core nodes of instance m4.large

General Configuration

Cluster name

☐ Logging ⓘ

Launch mode ☒ Cluster ⓘ ☐ Step execution ⓘ

Software configuration

Release ⓘ

Applications ☒ Core Hadoop: Hadoop 2.8.5 with Ganglia 3.7.2, Hive 2.3.6, Hue 4.4.0, Mahout 0.13.0, Pig 0.17.0, and Tez 0.9.2

☐ HBase: HBase 1.4.10 with Ganglia 3.7.2, Hadoop 2.8.5, Hive 2.3.6, Hue 4.4.0, Phoenix 4.14.3, and ZooKeeper 3.4.14

☐ Presto: Presto 0.227 with Hadoop 2.8.5 HDFS and Hive 2.3.6 Metastore

☐ Spark: Spark 2.4.4 on Hadoop 2.8.5 YARN with Ganglia 3.7.2 and Zeppelin 0.8.2

☐ Use AWS Glue Data Catalog for table metadata ⓘ

Hardware configuration

Instance type ⓘ The selected instance type adds 32 GiB of GP2 EBS storage per instance by default. [Learn more](#) ⓘ

Number of instances (1 master and 1 core nodes)

Security and access

EC2 key pair ⓘ [Learn how to create an EC2 key pair.](#)

Permissions ☒ Default ☐ Custom

Use default IAM roles. If roles are not present, they will be automatically created for you with managed policies for automatic policy updates.

EMR role [EMR_DefaultRole](#) ⓘ ☐ Use EMR_DefaultRole_V2 ⓘ

EC2 instance profile [EMR_EC2_DefaultRole](#) ⓘ

Cancel

Create cluster

Following that, YARN Configuration was done on the cluster

Clone Terminate AWS CLI export ▲ Auto-termination is not available for this account when using this release of EMR.

Cluster: upgrad-case study Running Running step

Summary Application user interfaces Monitoring Hardware Configurations Events Steps Bootstrap actions

Configuration classifications allow you to customize parameters for cluster applications. Instance groups inherit cluster configurations that are specified when a cluster is created. With EMR 5.21.0 and later, you can override cluster configurations and specify additional configuration classifications for each instance group below. [Learn more](#)

Cancel editing Save changes Instance group: ig-1G0B9V4PPEW61 Running

Classification	Property	Value	Source
yarn-site	yarn.scheduler.maximum-allocation-mb	8192	Instance Group configuration ✕
yarn-site	yarn.nodemanager.resource.memory-mb	10240	Instance Group configuration ✕

+ Add configuration

☒ Apply this configuration to all active instance groups

Moving data from S3 bucket into the HDFS

```
[hadoop@ip-172-31-18-85 ~]$ aws s3 cp s3://tutorialprashantk/2019-Nov.csv .
download: s3://tutorialprashantk/2019-Nov.csv to ./2019-Nov.csv
[hadoop@ip-172-31-18-85 ~]$ aws s3 cp s3://tutorialprashantk/2019-Oct.csv .
download: s3://tutorialprashantk/2019-Oct.csv to ./2019-Oct.csv
```

Tables clickstream_nov and clickstream_oct are created using CSVSerde and the data is loaded from CSV files in HDFS

```
hive> CREATE EXTERNAL TABLE clickstream_nov(
  > event_time timestamp,
  > event_type string,
  > product_id string,
  > category_id string,
  > category_code string,
  > brand string,
  > price float,
  > user_id bigint,
  > user_session string)
  > ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde'
  > WITH SERDEPROPERTIES (
  >   "separatorChar" = ",",
  > )
  > STORED AS TEXTFILE
  > TBLPROPERTIES ("skip.header.line.count"="1");
OK
Time taken: 0.662 seconds
hive>
  > CREATE EXTERNAL TABLE clickstream_oct(
  > event_time timestamp,
  > event_type string,
  > product_id string,
  > category_id string,
  > category_code string,
  > brand string,
  > price float,
  > user_id bigint,
  > user_session string)
  > ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde'
  > WITH SERDEPROPERTIES (
  >   "separatorChar" = ",",
  > )
  > STORED AS TEXTFILE
  > TBLPROPERTIES ("skip.header.line.count"="1");
OK
Time taken: 0.101 seconds
hive> load data local inpath '/home/hadoop/2019-Nov.csv' into table clickstream_nov;
Loading data to table upgrad.clickstream_nov
OK
Time taken: 9.568 seconds
hive> load data local inpath '/home/hadoop/2019-Oct.csv' into table clickstream_oct;
Loading data to table upgrad.clickstream_oct
OK
Time taken: 7.803 seconds
```

Using optimized techniques to run queries efficiently

Enabling Dynamic Partitioning

```
hive> set hive.exec.dynamic.partition = true;
hive> set hive.exec.dynamic.partition.mode = nonstrict;
```

Creating a partitioned and a bucketed table 'clickstream' to put all the data together

```
hive> CREATE EXTERNAL TABLE clickstream(
  > event_time timestamp,
  > product_id string,
  > category_id string,
  > category_code string,
  > brand string,
  > price float,
  > user_id bigint,
  > user_session string)
  > PARTITIONED BY (event_type string)
  > CLUSTERED by (category_id) into 7 buckets
  > ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde'
  > WITH SERDEPROPERTIES (
  >   "separatorChar" = ",",
  > )
  > stored as textfile;
OK
Time taken: 0.1 seconds
```

Loading data into 'clickstream' table

```
Time taken: 0.1 seconds
hive> insert into table clickstream partition (event_type) select event_time, product_id, category_id, category_code, brand, price,
  user_id, user_session, event_type from clickstream_nov;
Query ID = hadoop_20221130153457_91dc2a31-15d9-4473-91d0-2819f427ded8
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1669820927340_0003)

-----
VERTICES    MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container    SUCCEEDED    5         5         0         0         0         0
Reducer 2 ..... container    SUCCEEDED    3         3         0         0         0         0
-----
VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 106.96 s
-----
Loading data to table upgrad.clickstream partition (event_type=null)

Loaded : 5/5 partitions.
Time taken to load dynamic partitions: 0.363 seconds
Time taken for adding to write entity : 0.001 seconds
OK
Time taken: 108.476 seconds
hive> insert into table clickstream partition (event_type) select event_time, product_id, category_id, category_code, brand, price,
  user_id, user_session, event_type from clickstream_oct;
Query ID = hadoop_20221130153645_c812ae09-f34c-4a69-98b8-987671f8655c
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1669820927340_0003)

-----
VERTICES    MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container    SUCCEEDED    8         8         0         0         0         0
Reducer 2 ..... container    SUCCEEDED    5         5         0         0         0         0
-----
VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 137.01 s
-----
Loading data to table upgrad.clickstream partition (event_type=null)

Loaded : 5/5 partitions.
Time taken to load dynamic partitions: 0.482 seconds
Time taken for adding to write entity : 0.001 seconds
OK
Time taken: 138.786 seconds
```

Improvement of performance after using optimization

- Query for getting total revenue in the month of October was launched on clickstream, our partitioned and bucketed table, and clickstream_oct, the table we used to load data from csv files.
- Query on the clickstream table took 29.54 seconds whereas query on clickstream_oct table took 64.1 seconds.

```
hive> select sum(price) as total_revenue from clickstream where event_type = 'purchase' and month(event_time) = 10;
Query ID = hadoop_20221130153928_2cf8fa39-0490-4ea7-afdc-57216ba9c559
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1669820927340_0003)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    4         4         0         0         0         0
Reducer 2 ..... container  SUCCEEDED    1         1         0         0         0         0
-----
VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 28.92 s
-----
OK
1211538.4299999191
Time taken: 29.549 seconds, Fetched: 1 row(s)
hive> select sum(price) as total_revenue from clickstream_oct where event_type = 'purchase';
Query ID = hadoop_20221130154031_f63af778-64cd-4f11-8b6d-7ca38408fb38
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1669820927340_0003)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    8         8         0         0         0         0
Reducer 2 ..... container  SUCCEEDED    1         1         0         0         0         0
-----
VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 63.62 s
-----
OK
2742555.329999948
Time taken: 64.141 seconds, Fetched: 1 row(s)
```

Analysis – Answering Questions using HIVE Queries

Find the total revenue generated due to purchases made in October.

Answer: 12,11,538.43

```
hive> select sum(price) as total_revenue from clickstream where event_type = 'purchase' and month(event_time) = 10;
Query ID = hadoop_20221130153928_2cf8fa39-0490-4ea7-afdc-57216ba9c559
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1669820927340_0003)

-----
VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED   4        4          0        0        0        0
Reducer 2 ..... container  SUCCEEDED   1        1          0        0        0        0
-----
VERTICES: 02/02 [=====>>] 100% ELAPSED TIME: 28.92 s
-----
OK
1211538.4299999191
Time taken: 29.549 seconds, Fetched: 1 row(s)
```

Write a query to yield the total sum of purchases per month in a single output.

```
hive> select month(event_time) as month, sum(price) as sum_of_purchases from clickstream where event_type = 'purchase' group by month(event_time);
Query ID = hadoop_20221130154501_0112a231-bbd5-4e63-b047-f69703b8400d
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1669820927340_0003)

-----
VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED   4        4          0        0        0        0
Reducer 2 ..... container  SUCCEEDED   1        1          0        0        0        0
-----
VERTICES: 02/02 [=====>>] 100% ELAPSED TIME: 30.30 s
-----
OK
10      1211538.4299999191
11      3062033.7999994755
Time taken: 30.927 seconds, Fetched: 2 row(s)
```

Write a query to find the change in revenue generated due to purchases from October to November.

Answer: 1850495.37

```
hive> with monthwiserevenue as(
  > select month(event_time) as month, sum(price) as revenue from clickstream where event_type = 'purchase' group by month(event_time))
  > select revenue - lag(revenue) over (order by month) as change_in_revenue from monthwiserevenue;
Query ID = hadoop_20221130160142_66f52e8e-22a4-4aec-98b0-cb1c961d03e5
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1669820927340_0004)

-----
VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED   4        4          0        0        0        0
Reducer 2 ..... container  SUCCEEDED   1        1          0        0        0        0
Reducer 3 ..... container  SUCCEEDED   1        1          0        0        0        0
-----
VERTICES: 03/03 [=====>>] 100% ELAPSED TIME: 31.33 s
-----
OK
NULL
1850495.3699995563
Time taken: 32.173 seconds, Fetched: 2 row(s)
```

Find distinct categories of products. Categories with null category code can be ignored.

```
hive> select distinct category_code from clickstream where category_code is not null;
Query ID = hadoop_20221130171842_139b8fd4-71a8-40ea-9cf2-5580194ae0bb
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1669827995695_0001)

-----
VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED   6         6         0         0         0         0
Reducer 2 ..... container  SUCCEEDED   5         5         0         0         0         0
-----
VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 70.96 s
-----
OK
category_code

accessories.cosmetic_bag
stationery.cartridge
accessories.bag
appliances.environment.vacuum
furniture.living_room.chair
sport.diving
appliances.personal.hair_cutter
appliances.environment.air_conditioner
apparel.glove
furniture.bathroom.bath
furniture.living_room.cabinet
Time taken: 72.087 seconds, Fetched: 12 row(s)
```

Find the total number of products available under each category.

```
hive> select category_code, count(product_id) as number_of_products from clickstream where category_code is not null group by category_code;
Query ID = hadoop_20221130172759_80506770-e076-4fef-8711-ec89fa486b18
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1669827995695_0001)

-----
VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED   6         6         0         0         0         0
Reducer 2 ..... container  SUCCEEDED   5         5         0         0         0         0
-----
VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 70.62 s
-----
OK
category_code  number_of_products
8594895
accessories.cosmetic_bag      1248
stationery.cartridge          26722
accessories.bag 11681
appliances.environment.vacuum  59761
furniture.living_room.chair    308
sport.diving 2
appliances.personal.hair_cutter 1643
appliances.environment.air_conditioner 332
apparel.glove 18232
furniture.bathroom.bath 9857
furniture.living_room.cabinet 13439
Time taken: 71.321 seconds, Fetched: 12 row(s)
```

Which brand had the maximum sales in October and November combined?

Answer: runail

```
hive> select brand, sum(price) as sales from clickstream where brand is not null and brand != '' and event_type = 'purchase' group by brand order by sales desc limit 1;
Query ID = hadoop_20221130174646_668faded-7309-45d7-8ffe-953985de5e05
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1669827995695_0001)

-----
VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED   3         3         0         0         0         0
Reducer 2 ..... container  SUCCEEDED   1         1         0         0         0         0
Reducer 3 ..... container  SUCCEEDED   1         1         0         0         0         0
-----
VERTICES: 03/03 [=====>>>] 100% ELAPSED TIME: 24.49 s
-----
OK
brand  sales
runail 148297.93999998464
Time taken: 25.349 seconds, Fetched: 1 row(s)
```

Which brands increased their sales from October to November?

There are a total of 152 brands that increased sales from October to November.

```
hive> with brands_sales_oct as (  
> select brand, sum(price) as oct_sales from clickstream where event_type = 'purchase' and brand is not null and brand != '' and month(event_time) = 10 group by b  
rand  
> ),  
> brands_sales_nov as (  
> select brand, sum(price) as nov_sales from clickstream where event_type = 'purchase' and brand is not null and brand != '' and month(event_time) = 11 group by b  
rand  
> )  
> select o.brand as brand, n.nov_sales - o.oct_sales as increase_in_sales from brands_sales_oct o inner join brands_sales_nov n on o.brand = n.brand where n.nov_sa  
les - o.oct_sales > 0 order by increase_in_sales desc;  
Query ID = hadoop_20221130180405_4448e44d-0ffa-4a2e-9482-76a7d21aad0e  
Total jobs = 1  
Launching Job 1 out of 1  
Status: Running (Executing on YARN cluster with App id application_1669827995695_0001)
```

VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	container	SUCCEEDED	3	3	0	0	0	0
Map 5	container	SUCCEEDED	3	3	0	0	0	0
Reducer 2	container	SUCCEEDED	1	1	0	0	0	0
Reducer 3	container	SUCCEEDED	1	1	0	0	0	0
Reducer 4	container	SUCCEEDED	1	1	0	0	0	0
Reducer 6	container	SUCCEEDED	1	1	0	0	0	0

VERTICES: 06/06 [=====] 100% ELAPSED TIME: 34.62 s

```
OK  
brand      increase_in_sales  
grattol    36027.16999999717  
uno        15737.71999999987  
lianail    10501.400000000256  
ingarden   10404.819999998246  
strong     9474.639999999941  
jessnail   7057.389999999977  
cosmoengi  5214.180000000126  
polarus    5358.2100000000055  
runail     5219.3799999987461  
freedecor  4250.020000000293  
staleks    3355.8799999999537  
bpw.style  3265.2900000009977  
lovely     3234.679999999991  
marathon   2992.350000000003  
haruyama   2562.2199999996046  
yoko       2950.970000000103  
italwax    2859.1300000003175  
benovy     2850.3499999999967  
kaypro     2387.36  
estel      2385.9199999998928  
concept    2348.260000000053  
kapous     2165.920000000002  
f.o.x      1953.0499999999656  
masura     1792.39000000007234  
milv       1737.0700000000502  
beautix    1729.000000000011  
artex      1596.6099999999942  
domix      1537.119999999999  
shik       1498.5200000000023  
smart      1444.880000000021  
rouloff    1422.4899999999999  
levrana    1420.5400000000004  
onig       1416.2399999999798  
irisk      1354.07999999994633  
severina   1344.600000000004  
joico      1309.5799999999997  
zeitun     1300.9700000000007  
beauty-free 1228.6900000000132  
swarovski  1155.2300000000162  
de.lux     1115.809999999913  
metzger    1083.7099999999802  
markell    1065.6800000000012  
sanoto     1052.54  
nagaraku   957.9400000000007  
ecolab     951.4499999999999  
art-visage 905.0899999999992  
levissime  857.0099999999904  
missha     856.4500000000003  
solomeya   786.0099999999938  
rosi       764.5200000000015  
refectocil 759.4000000000042  
kaaral     673.6399999999894  
kosmekka   631.9300000000005  
kinetics   611.0099999999829  
browxenna  585.3600000000006  
airmails   572.6200000000581  
uskul      548.0400000000345  
coifin     525.4900000000001  
s.care     500.3899999999993  
limoni     487.70000000000095  
matrix     483.4899999999989  
gehwol     468.6099999999997  
greymy     460.2800000000003  
bioaqua    455.2299999999997  
farmavita  454.60000000000036  
sophin     447.6599999999985  
yu-r       402.2999999999967  
kiss       395.7800000000001  
lador      387.91999999999234  
ellips     360.1899999999999  
jas        338.4699999999935  
lowence    324.9100000000001  
nitrile    315.40000000000043  
shary      304.5299999999986  
kims       302.0000000000006  
happyfons  289.6699999999985  
kocostar   284.0000000000001  
insight    278.2599999999954  
candy      264.4200000000003  
bluesky    258.28999999998608  
beaugreen  256.84000000000003  
protokeratin 255.54000000000008  
trind      244.89000000000004  
entity     239.54999999999808  
skinlite   238.51000000000056  
provoc     235.83000000000243  
fedua      211.43  
ecocraft   200.79  
keen       199.2699999999998  
mane       193.47  
freshbubble 183.64  
chi        179.67000000000002  
cristalinas 157.31999999999988  
farmona    150.9700000000007
```



```

farmona 150.9700000000007
latinoil 135.07000000000005
miskin 135.03
elizavecca 133.77
nefertiti 133.11999999999992
finish 132.0
igroebauty 131.409999999999917
dizao 126.30000000000042
osmo 116.72999999999999
batiiste 101.76999999999975
carmex 98.28
eos 98.27000000000004
depilflax 96.710000000000322
enjoy 95.22000000000003
kerasys 94.29000000000008
aura 93.55999999999996
plazan 92.64000000000001
koelf 84.56000000000034
nirvel 71.29000000000013
konad 70.84000000000026
egomania 68.57
cutrin 68.25
laboratorium 66.0200000000001
imn 63.19000000000028
marutaka-foot 60.110000000000014
profhenna 57.61999999999966
koelcia 57.25
balbcare 57.05000000000024
elskin 56.56000000000003
foamie 45.44999999999998
ladykin 44.92000000000003
likato 44.91000000000025
mavala 37.27999999999997
vilenta 33.60999999999985
beautyblender 30.66999999999987
biore 29.65999999999997
orly 28.70999999999958
estelare 27.060000000000343
profepil 24.65999999999997
blixz 24.44999999999999
godefroy 23.899999999999864
glysolid 21.85999999999997
veraclara 21.10000000000003
kamill 19.40000000000004
treaclemoon 18.120000000000005
supertan 16.140000000000008
deoproce 12.32999999999927
rasyan 10.13999999999997
fly 10.030000000000001
tertio 9.640000000000001
jaguar 8.53999999999964
soleo 8.329999999999671
neoleor 8.290000000000006
moyou 4.570000000000001
bodyton 4.299999999999272
skinity 3.560000000000005
grace 1.690000000000012
cosima 0.700000000000028
ovale 0.56
Time taken: 36.015 seconds, Fetched: 152 row(s)

```

Write a query to generate a list of top 10 users who spend the most.

```

hive> select user_id, sum(price) as total_purchases from clickstream where event_type = 'purchase' and user_id is not null group by user_id order by total_purchases desc limit 10;
Query ID = hadoop_20221130181438_316b1f2e-98d2-4540-9724-231a320b481d
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1669827995695_0001)

-----
VERTICES      MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED  3      3      0      0      0      0
Reducer 2 ..... container  SUCCEEDED  1      1      0      0      0      0
Reducer 3 ..... container  SUCCEEDED  1      1      0      0      0      0
-----
VERTICES: 03/03 [=====] 100% ELAPSED TIME: 28.69 s
-----
OK
user_id total_purchases
557790271 2715.8699999999994
150318419 1645.9699999999998
562167663 1352.85
531900924 1329.4499999999996
557850743 1295.4799999999998
522130011 1185.3900000000001
561592095 1109.7000000000001
431950134 1097.59
566576008 1056.3600000000013
521347209 1040.9099999999999
Time taken: 29.268 seconds, Fetched: 10 row(s)

```

Cleaning up

Dropping the Database

```
hive> show tables;
OK
tab_name
clickstream
clickstream_nov
clickstream_oct
Time taken: 0.036 seconds, Fetched: 3 row(s)
hive> drop table clickstream;
OK
Time taken: 0.195 seconds
hive> drop table clickstream_nov;
OK
Time taken: 0.123 seconds
hive> drop table clickstream_oct;
OK
Time taken: 0.071 seconds
hive> drop database upgrad;
OK
Time taken: 0.196 seconds
hive> show databases;
OK
database_name
default
Time taken: 0.011 seconds, Fetched: 1 row(s)
```

Terminating the Cluster

Clone

Terminate

AWS CLI export

⚠️ Auto-termination is not available for this account when using this release of EMR.

Cluster: upgrad-case study Terminating Terminated by user request

Summary

Application user interfaces

Monitoring

Hardware

Configurations

Events

Steps

Bootstrap actions

Summary

ID: j-BH0FED97ONFD
Creation date: 2022-11-30 22:28 (UTC+5:30)
Elapsed time: 1 hour, 28 minutes
After last step completes: Cluster waits
Termination protection: Off
Tags: --
Master public DNS: ec2-54-147-187-196.compute-1.amazonaws.com [🔗](#)
Connect to the Master Node Using SSH

Configuration details

Release label: emr-5.29.0
Hadoop distribution: Amazon 2.8.5
Applications: Ganglia 3.7.2, Hive 2.3.6, Hue 4.4.0, Mahout 0.13.0, Pig 0.17.0, Tez 0.9.2
Log URI: --
EMRFS consistent view: Disabled
Custom AMI ID: --

Application user interfaces

Persistent user interfaces [🔗](#): --
On-cluster user -- interfaces [🔗](#):

Network and hardware

Availability zone: us-east-1d
Subnet ID: [subnet-0ade01654002175da](#) [🔗](#)
Master: Running 1 m4.large
Core: Running 1 m4.large
Task: --
Cluster scaling: Not enabled

Security and access

Key name: vockey
EC2 instance profile: EMR_EC2_DefaultRole
EMR role: EMR_DefaultRole
Visible to all users: All [Change](#)
Security groups for Master: [sg-0541e5e31b09511e8](#) [🔗](#) (ElasticMapReduce-master)
Security groups for Core & Task: [sg-0bce01855743c317b](#) [🔗](#) (ElasticMapReduce-slave)

Appendix

Script of all the commands used:

#Move Data from S3 to HDFS

```
aws s3 cp s3://tutorialprashantk/2019-Nov.csv .
```

```
aws s3 cp s3://tutorialprashantk/2019-Oct.csv .
```

#Creating a database on Hive

```
create database upgrad;
```

```
use upgrad;
```

#Printing Headers

```
set hive.cli.print.header=true ;
```

#Creating a table

```
CREATE EXTERNAL TABLE clickstream_nov(
```

```
event_time timestamp,
```

```
event_type string,
```

```
product_id string,
```

```
category_id string,
```

```
category_code string,
```

```
brand string,
```

```
price float,
```

```
user_id bigint,
```

```
user_session string)
```

```
ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde'
```

```
WITH SERDEPROPERTIES (
```

```
    "separatorChar" = ",",
```

```
)
```

```
STORED AS TEXTFILE
```

```
TBLPROPERTIES ("skip.header.line.count"="1");
```

```
CREATE EXTERNAL TABLE clickstream_oct(  
  event_time timestamp,  
  event_type string,  
  product_id string,  
  category_id string,  
  category_code string,  
  brand string,  
  price float,  
  user_id bigint,  
  user_session string)  
ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde'  
WITH SERDEPROPERTIES (  
  "separatorChar" = ","  
)  
STORED AS TEXTFILE  
TBLPROPERTIES ("skip.header.line.count"="1");
```

#Loading Data

```
load data local inpath '/home/hadoop/2019-Nov.csv' into table clickstream_nov;  
load data local inpath '/home/hadoop/2019-Oct.csv' into table clickstream_oct;
```

#Enabling Dynamic Partitioning

```
set hive.exec.dynamic.partition = true;  
set hive.exec.dynamic.partition.mode = nonstrict;
```

#Creating our final table with partitions and buckets

```
CREATE EXTERNAL TABLE clickstream(  
  event_time timestamp,  
  product_id string,  
  category_id string,  
  category_code string,  
  brand string,  
  price float,  
  user_id bigint,  
  user_session string)  
PARTITIONED BY (event_type string)  
CLUSTERED by (category_id) into 7 buckets  
ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde'  
WITH SERDEPROPERTIES (  
  "separatorChar" = ",",  
)  
stored as textfile  
TBLPROPERTIES ('serialization.null.format'='');
```

#Inserting data

```
insert into table clickstream partition (event_type) select event_time, product_id, category_id,  
category_code, brand, price, user_id, user_session, event_type from clickstream_nov;
```

```
insert into table clickstream partition (event_type) select event_time, product_id, category_id,  
category_code, brand, price, user_id, user_session, event_type from clickstream_oct;
```

----- Getting Answers for questions -----

#Find the total revenue generated due to purchases made in October.

```
select sum(price) as total_revenue from clickstream where event_type = 'purchase' and  
month(event_time) = 10;
```

#Write a query to yield the total sum of purchases per month in a single output.

```
select month(event_time) as month, sum(price) as sum_of_purchases from clickstream where  
event_type = 'purchase' group by month(event_time);
```

#Write a query to find the change in revenue generated due to purchases from October to November.

with monthwiserevenue as(

```
select month(event_time) as month, sum(price) as revenue from clickstream where event_type =  
'purchase' group by month(event_time))
```

```
select revenue - lag(revenue) over (order by month) as change_in_revenue from monthwiserevenue;
```

#Find distinct categories of products. Categories with null category code can be ignored
select distinct category_code from clickstream where category_code is not null;

#Find the total number of products available under each category.
select category_code, count(product_id) as number_of_products from clickstream where category_code is not null group by category_code;

#Which brand had the maximum sales in October and November combined?
select brand, sum(price) as sales from clickstream where brand is not null and brand != "" and event_type = 'purchase' group by brand order by sales desc limit 1;

#Which brands increased their sales from October to November?
with brands_sales_oct as (
select brand, sum(price) as oct_sales from clickstream where event_type = 'purchase' and brand is not null and brand != "" and month(event_time) = 10 group by brand
)
brands_sales_nov as (
select brand, sum(price) as nov_sales from clickstream where event_type = 'purchase' and brand is not null and brand != "" and month(event_time) = 11 group by brand
)
select o.brand as brand, n.nov_sales - o.oct_sales as increase_in_sales from brands_sales_oct o inner join brands_sales_nov n on o.brand = n.brand where n.nov_sales - o.oct_sales > 0 order by increase_in_sales desc;

#Your company wants to reward the top 10 users of its website with a Golden Customer plan. Write a query to generate a list of top 10 users who spend the most.

select user_id, sum(price) as total_purchases from clickstream where event_type = 'purchase' and user_id is not null group by user_id order by total_purchases desc limit 10;